

final_recovery_time

2023-05-10

Data Cleaning

```
dat0 = read_csv(file = "final_used_data.csv")

## Rows: 3604 Columns: 13
## -- Column specification -----
## Delimiter: ","
## dbl (13): id, age, gender, race, smoking, bmi, hypertension, diabetes, sbp, ...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.


dat1 <- na.omit(dat0)

#Create a new data frame without id variable
dat <- dat1[ , !names(dat0) %in% c("id")]
attach(dat)
```


Data Preprocessing

```
#Visualize response variable
par(mfrow = c(1, 2))
boxplot(dat$recovery_time, main = "COVID-19 Recovery Time")
hist(dat$recovery_time, main = "Distribution of Rescovery Time", col = "lightblue",
      xlab = "Recovery Time", prob = TRUE, ylim = c(0,0.03))
lines(density(dat$recovery_time))
abline(v = mean(dat$recovery_time), lty = "dashed", col = "red")
```

COVID-19 Recovery Time




Distribution of Recovery Time




The above plots show that the response variable is right-skewed, so log-transformation was performed.

```
dat$log_recovery_time <- log(dat$recovery_time)
par(mfrow = c(1, 2))
boxplot(dat$log_recovery_time, main = "COVID-19 Recovery Time")
hist(dat$log_recovery_time, main = "Distribution of Recovery Time", col = "lightblue",
xlab = "Log Recovery Time", prob = TRUE, ylim = c(0,1))
lines(density(dat$log_recovery_time))
abline(v = mean(dat$log_recovery_time), lty = "dashed", col = "red")
```

COVID-19 Recovery Time



Distribution of Recovery Time



#After the transformation, the outcome variable is normally distributed.

Data Partition

```
set.seed(5220)
trainRows <- createDataPartition(y = dat$log_recovery_time, p = 0.8, list = FALSE)

# Training data
dat_train = dat[trainRows, ]
x_train = model.matrix(log_recovery_time~.,dat)[trainRows, -1]
y_train = dat$log_recovery_time[trainRows]

# Test data
dat_test = dat[-trainRows, ]
x_test = model.matrix(log_recovery_time~.,dat)[-trainRows, -1]
y_test = dat$log_recovery_time[-trainRows]
```

Exploratory Analysis & Data Visualization

```
# Summary statistics for each variable
summary(dat_train)
```

```
##      age      gender      race      smoking
## Min.   :46.00  Min.   :0.0000  Min.   :1.000  Min.   :0.0000
## 1st Qu.:57.00  1st Qu.:0.0000  1st Qu.:1.000  1st Qu.:0.0000
```


```
## Median :60.00 Median :0.0000 Median :1.000 Median :0.0000
## Mean :60.11 Mean :0.4884 Mean :1.753 Mean :0.4919
## 3rd Qu.:63.00 3rd Qu.:1.0000 3rd Qu.:3.000 3rd Qu.:1.0000
## Max. :79.00 Max. :1.0000 Max. :4.000 Max. :2.0000
##      bmi      hypertension      diabetes      sbp
## Min. :19.7 Min. :0.0000 Min. :0.0000 Min. :103.0
## 1st Qu.:25.8 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:125.0
## Median :27.6 Median :0.0000 Median :0.0000 Median :130.0
## Mean :27.7 Mean :0.4853 Mean :0.1622 Mean :130.3
## 3rd Qu.:29.5 3rd Qu.:1.0000 3rd Qu.:0.0000 3rd Qu.:136.0
## Max. :39.8 Max. :1.0000 Max. :1.0000 Max. :157.0
##      ldl      vaccine      severity      recovery_time
## Min. : 45.0 Min. :0.0000 Min. :0.00000 Min. : 3.00
## 1st Qu.: 97.0 1st Qu.:0.0000 1st Qu.:0.00000 1st Qu.: 28.00
## Median :110.0 Median :1.0000 Median :0.00000 Median : 39.00
## Mean :110.1 Mean :0.5934 Mean :0.09185 Mean : 43.38
## 3rd Qu.:124.0 3rd Qu.:1.0000 3rd Qu.:0.00000 3rd Qu.: 50.00
## Max. :174.0 Max. :1.0000 Max. :1.00000 Max. :365.00
## log_recovery_time
## Min. :1.099
## 1st Qu.:3.332
## Median :3.664
## Mean :3.612
## 3rd Qu.:3.912
## Max. :5.900
```

```
#Relocate columns putting non-discrete predictors together
```

```
dat1 =
  dat %>%
  select(age,bmi,sbp,ldl,log_recovery_time)
```

```
# Correlation Plot
```

```
dat2 <- model.matrix(log_recovery_time ~ ., dat1)[ , -1]
x <- dat2[trainRows,]
corrplot(cor(x))
```



```
# Convert non-numeric columns to numeric
dat_train1 <- dat_train
non_numeric_cols <- sapply(dat_train1, function(x) !is.numeric(x))
dat_train1[, non_numeric_cols] <- lapply(dat_train1[, non_numeric_cols], as.numeric)
```


Fit LASSO

```
ctrl1 <- trainControl(method = "repeatedcv", number = 10, repeats = 5)

set.seed(5220)
lasso.fit <- train(x_train, y_train,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = 1,
    lambda = exp(seq(5, -1, length=100))),
  trControl = ctrl1)
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.
```

```
plot(lasso.fit, xTrans = log)
```



```
lasso.fit$bestTune
```

```
##   alpha   lambda
## 1      1 0.3678794
```

```
coef(lasso.fit$finalModel, lasso.fit$bestTune$lambda)
```

```
## 13 x 1 sparse Matrix of class "dgCMatrix"
##               s1
## (Intercept)  3.454454838
## age          .
## gender       .
## race         .
## smoking      .
## bmi          .
## hypertension .
## diabetes     .
## sbp          .
## ldl          .
## vaccine      .
## severity     .
## recovery_time 0.003622746
```

Fit Elastic net

```

set.seed(5220)
enet.fit <- train(x_train, y_train,
                 method = "glmnet",
                 tuneGrid = expand.grid(alpha = seq(0, 1, length = 21),
                 lambda = exp(seq(2, -2, length = 50))),
                 trControl = ctrl1)

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

```

```
enet.fit$bestTune
```

```

##   alpha    lambda
## 1      0 0.1353353

```


```
coef(enet.fit$finalModel, enet.fit$bestTune$lambda)
```

```

## 13 x 1 sparse Matrix of class "dgCMatrix"
##               s1
## (Intercept)  2.6435219747
## age          0.0018941343
## gender       -0.0402957253
## race         -0.0054230774
## smoking      0.0273289815
## bmi          0.0031212386
## hypertension 0.0120551948
## diabetes     0.0095582923
## sbp          0.0018610716
## ldl          0.0001388527
## vaccine     -0.0583915581
## severity     0.0563081990
## recovery_time 0.0126284871

```

```
plot(enet.fit)
```




Fit PCR

```
ctrl2 <- trainControl(method = "repeatedcv",
                      number = 10,
                      repeats = 5,
                      selectionFunction = "best")

set.seed(5220)
pcr.fit <- train(x_train, y_train,
                 method = "pcr",
                 tuneGrid = data.frame(ncomp = 1:13),
                 trControl = ctrl2,
                 preProcess = c("center", "scale"))

ggplot(pcr.fit, highlight = TRUE) + theme_bw()
```





```
summary(pcr.fit)
```

```
## Data:      X dimension: 2885 12
## Y dimension: 2885 1
## Fit method: svdpc
## Number of components considered: 11
## TRAINING: % variance explained
##           1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps
## X           18.536  30.03  38.94  47.60  55.99  64.29  72.29
## .outcome     2.576  46.70  46.82  47.14  47.17  47.32  47.34
##           8 comps 9 comps 10 comps 11 comps
## X           80.14  87.41  93.17  98.43
## .outcome     48.22  48.23  48.25  69.98
```

Fit PLS

```
set.seed(5220)
pls.fit <- train(x_train, y_train,
                 method = "pls",
                 tuneGrid = data.frame(ncomp = 1:13),
                 trControl = ctrl2,
                 preProcess = c("center", "scale"))
ggplot(pls.fit, highlight = TRUE)
```



```
summary(pls.fit)
```

```
## Data:      X dimension: 2885 12
## Y dimension: 2885 1
## Fit method: oscorespls
## Number of components considered: 4
## TRAINING: % variance explained
##           1 comps 2 comps 3 comps 4 comps
## X           11.78  25.05  35.33  43.41
## .outcome     62.66  68.25  69.97  69.99
```

Fit GAM

```
#ctrl3 <- trainControl(method = "cv", number = 10)

set.seed(5220)
gam.fit <- train(x_train, y_train,
                 method = "gam",
                 trControl = ctrl1)
```

```
## Loading required package: mgcv
```

```
## Loading required package: nlme
```

```
##
## Attaching package: 'nlme'

## The following object is masked from 'package:dplyr':
##
## collapse

## This is mgcv 1.8-40. For overview type 'help("mgcv-package")'.
```

```
gam.fit$bestTune
```

```
## select method
## 2 TRUE GCV.Cp
```

```
gam.fit$finalModel
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## .outcome ~ gender + hypertension + diabetes + vaccine + severity +
## smoking + race + s(age) + s(sbp) + s(ldl) + s(recovery_time) +
## s(bmi)
##
## Estimated degrees of freedom:
## 0.0004 0.0001 6.4058 8.9836 0.4052 total = 23.8
##
## GCV score: 0.003162374
```

```
coef(gam.fit$finalModel)
```

```
## (Intercept) gender hypertension diabetes
## 3.612330e+00 9.722452e-04 4.146820e-04 1.305287e-03
## vaccine severity smoking race
## -1.523296e-03 -4.759119e-03 2.982511e-04 -2.407263e-04
## s(age).1 s(age).2 s(age).3 s(age).4
## -2.718648e-09 -2.242637e-08 5.459337e-08 -8.260225e-08
## s(age).5 s(age).6 s(age).7 s(age).8
## 1.966523e-08 6.613925e-08 1.678210e-08 4.973789e-07
## s(age).9 s(sbp).1 s(sbp).2 s(sbp).3
## 2.112459e-10 -2.640486e-08 -4.580629e-09 1.930582e-08
## s(sbp).4 s(sbp).5 s(sbp).6 s(sbp).7
## -3.429016e-08 -1.392191e-08 2.878634e-08 1.522718e-08
## s(sbp).8 s(sbp).9 s(ldl).1 s(ldl).2
## 1.531279e-07 -1.093297e-10 1.401524e-03 2.314941e-02
## s(ldl).3 s(ldl).4 s(ldl).5 s(ldl).6
## -5.147928e-03 1.687779e-02 8.290092e-03 -8.597029e-03
## s(ldl).7 s(ldl).8 s(ldl).9 s(recovery_time).1
## -1.070984e-02 -5.783819e-02 -1.133501e-11 -1.101218e+00
## s(recovery_time).2 s(recovery_time).3 s(recovery_time).4 s(recovery_time).5
```


```
##      2.482169e+00      1.071460e+00      1.411544e+00      -9.561915e-01
## s(recovery_time).6 s(recovery_time).7 s(recovery_time).8 s(recovery_time).9
##      -1.022372e+00      1.227444e+00      1.282190e+00      1.173501e+00
##      s(bmi).1      s(bmi).2      s(bmi).3      s(bmi).4
##      -1.855887e-07      2.174533e-08      3.572935e-08      -4.140910e-08
##      s(bmi).5      s(bmi).6      s(bmi).7      s(bmi).8
##      -1.157483e-08      -3.957406e-08      7.307461e-09      2.533496e-07
##      s(bmi).9
##      5.648109e-04
```

```
mod_gam <- gam(log_recovery_time ~ gender + race + smoking + hypertension +
  diabetes + vaccine + severity + s(age) +
  s(sbp) + s(ldl) + s(bmi),
  data = dat[trainRows,], method = "REML")

summary(mod_gam)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## log_recovery_time ~ gender + race + smoking + hypertension +
##   diabetes + vaccine + severity + s(age) + s(sbp) + s(ldl) +
##   s(bmi)
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.728313   0.029602 125.947 < 2e-16 ***
## gender       -0.103779   0.018851  -5.505 4.02e-08 ***
## race        -0.016905   0.008700  -1.943  0.0521 .
## smoking      0.080160   0.013982   5.733 1.09e-08 ***
## hypertension 0.052285   0.033266   1.572  0.1161
## diabetes     0.005141   0.025571   0.201  0.8407
## vaccine     -0.197984   0.019165 -10.330 < 2e-16 ***
## severity     0.168279   0.032605   5.161 2.62e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F p-value
## s(age) 2.665  3.401  1.892  0.110
## s(sbp) 1.779  2.233  1.248  0.221
## s(ldl) 1.003  1.005  0.004  0.967
## s(bmi) 5.677  6.794 80.886 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.216   Deviance explained = 22.1%
## -REML =    2161   Scale est. = 0.25444   n = 2885
```

```
par(mfrow = c(1,6))
plot(gam.fit$finalModel)
```



Fit MARS

```

mars_grid <- expand.grid(degree = 1:3,
                        nprune = 2:15)
set.seed(5220)
mars.fit <- train(x_train, y_train,
                  method = "earth",
                  tuneGrid = mars_grid,
                  trControl = ctrl1)

```

```
## Loading required package: earth
```


```
## Loading required package: Formula
```

```
## Loading required package: plotmo
```

```
## Loading required package: plotrix
```

```
## Loading required package: TeachingDemos
```

```
ggplot(mars.fit)
```



```
mars.fit$bestTune
```

```
##      nprune degree
## 3         4       1
```

```
coef(mars.fit$finalModel)
```

```
##           (Intercept) h(recovery_time-41) h(41-recovery_time)
##           3.806773377          0.014172314          -0.043827302
## h(recovery_time-110)
##           -0.009387009
```


Model Comparison

```
set.seed(5220)
resamp <- resamples(list(enet =enet.fit, lasso = lasso.fit, pcr = pcr.fit, pls = pls.fit, gam=gam.fit))
summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models:  enet, lasso, pcr, pls, gam
## Number of resamples: 50
```

```
##
## MAE
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## enet  0.20448398 0.21367310 0.22406282 0.22380451 0.2316555 0.2466645    0
## lasso 0.31613346 0.34329906 0.35120802 0.35226982 0.3635228 0.3773298    0
## pcr   0.17894427 0.18824442 0.19810593 0.19799056 0.2062234 0.2174753    0
## pls   0.17894328 0.18816450 0.19819321 0.19800937 0.2062770 0.2176440    0
## gam   0.03090753 0.03313992 0.03433917 0.03432081 0.0353659 0.0400894    0
##
## RMSE
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## enet  0.26938541 0.29772082 0.32704779 0.32684611 0.34811402 0.38787385    0
## lasso 0.41847990 0.46740296 0.48034811 0.48332882 0.50435816 0.54279532    0
## pcr   0.23794846 0.28077381 0.31427071 0.31349502 0.33886646 0.39727401    0
## pls   0.23790094 0.28094331 0.31440206 0.31351276 0.33884139 0.39728622    0
## gam   0.03905939 0.04728467 0.05492136 0.05513242 0.06175933 0.08638066    0
##
## Rsquared
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## enet  0.6144698 0.6730869 0.7022847 0.7168417 0.7619367 0.8361868    0
## lasso 0.6072387 0.6688274 0.7144653 0.7218806 0.7724960 0.8578139    0
## pcr   0.6161576 0.6779777 0.7119874 0.7224607 0.7702780 0.8427547    0
## pls   0.6161506 0.6781600 0.7118646 0.7224473 0.7700829 0.8426681    0
## gam   0.9811190 0.9891391 0.9905767 0.9907163 0.9931621 0.9954404    0
```

```
bwplot(resamp, metric = "RMSE")
```



Regression tree

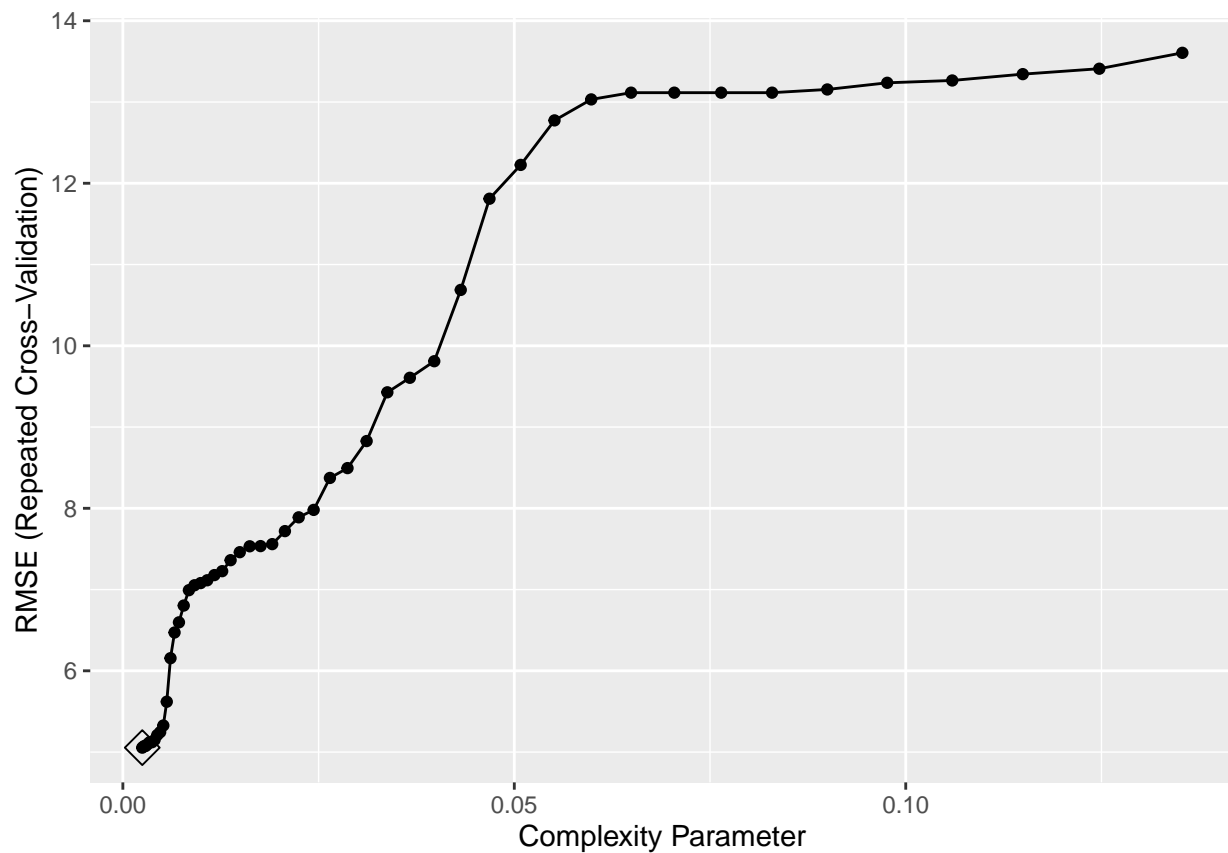
```
set.seed(5220)

rpart.fit <- train(recovery_time ~ . ,
                  data = dat_train,
                  method = "rpart",
                  tuneGrid = data.frame(cp = exp(seq(-6,-2, length = 50))),
                  trControl = ctrl1)


rpart.fit$bestTune
```

```
##           cp
## 1 0.002478752
```

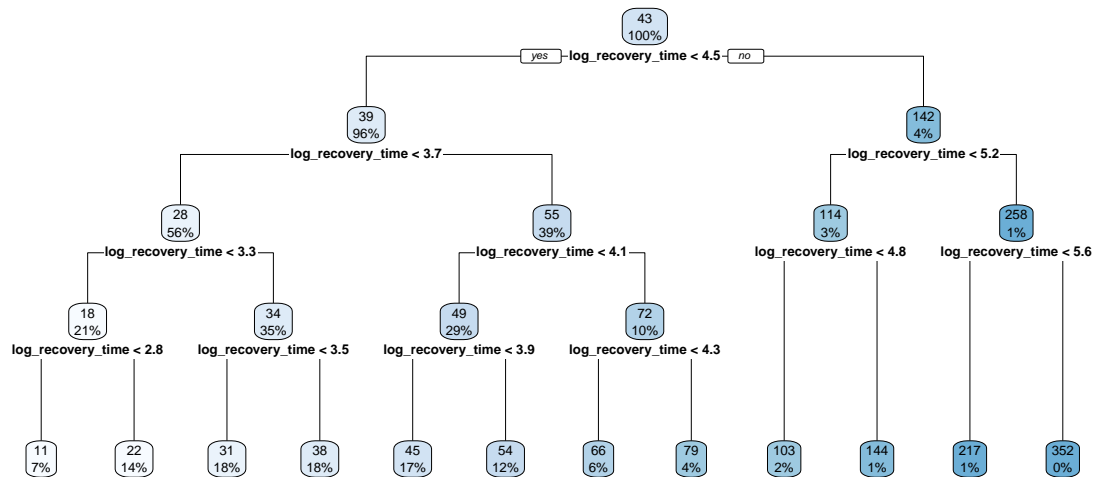
```
# Plot of the complexity parameter
ggplot(rpart.fit, highlight = TRUE)
```



```
# Variable importance
plot(varImp(rpart.fit, scale = TRUE))
```

```
rpart.plot(rpart.fit$finalModel)
```



```

pred.rf <- predict(rpart.fit, newdata = dat_test)
RMSE(pred.rf, dat_test$recovery_time)

```

```
## [1] 4.298444
```

Random forest

```


rf.grid <- expand.grid(mtry = 1:11,
                      splitrule = "variance",
                      min.node.size = 1:6)

set.seed(5220)

rf.fit <- train(recovery_time ~ .,
                data = dat_train,
                method = "ranger",
                tuneGrid = rf.grid,
                trControl = ctrl1)

ggplot(rf.fit, highlight = TRUE)

```




```
rf.fit$bestTune
```

```
##      mtry splitrule min.node.size
## 63    11  variance              3
```


```
# variable importance using permutation methods
set.seed(5220)
rf.perm = ranger(recovery_time ~ .,
                  data = dat_train,
                  mtry = rf.fit$bestTune[[1]],
                  splitrule = "variance",
                  min.node.size = rf.fit$bestTune[[3]],
                  importance = "permutation",
                  scale.permutation.importance = TRUE)

barplot(sort(ranger::importance(rf.perm), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("cyan", "blue"))(19))
```



```
# variable importance using impurity methods
set.seed(5220)
rf.impu <- ranger(recovery_time ~ .,
                  data = dat_train,
                  mtry = rf.fit$bestTune[[1]],
                  splitrule = "variance",
                  min.node.size = rf.fit$bestTune[[3]],
                  importance = "impurity")

barplot(sort(ranger::importance(rf.impu), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("cyan", "blue"))(19))
```



```
# test error
pred.rf <- predict(rf.fit, newdata = dat_test)
RMSE(pred.rf, dat_test$recovery_time)
```


```
## [1] 1.278566
```

Boosting

```
bst.grid = expand.grid(n.trees = c(2000,3000,4000,5000),
                      interaction.depth = 1:3,
                      shrinkage = c(0.005,0.01),
                      n.minobsinnode = c(1))

set.seed(5220)
bst.fit <- train(recovery_time ~.,
                 data = dat_train,
                 method = "gbm",
                 tuneGrid = bst.grid,
                 trControl = ctrl1,
                 verbose = FALSE)


ggplot(bst.fit, highlight = TRUE)
```



```
bst.fit$bestTune
```

```
##   n.trees interaction.depth shrinkage n.minobsinnode
## 8    5000                2    0.005                1
```

```
# variable importance
summary(bst.fit$finalModel, las = 2, cBars = 11, cex.names = 0.6)
```



```
##                                var      rel.inf
## log_recovery_time log_recovery_time 9.999825e+01
## bmi                bmi 5.949987e-04
## age                age 5.065697e-04
## ldl                ldl 3.628965e-04
## sbp                sbp 1.217016e-04
## hypertension      hypertension 7.082340e-05
## smoking            smoking 4.067221e-05
## race              race 1.696229e-05
## severity           severity 1.588669e-05
## gender             gender 1.236924e-05
## vaccine            vaccine 4.149641e-06
## diabetes           diabetes 0.000000e+00
```

```
# test error
```

```
pred.bst <- predict(bst.fit, newdata = dat_test)
RMSE(pred.bst, dat_test$recovery_time)
```

```
## [1] 0.8662751
```

Model Comparison

```
set.seed(5220)
resamp <- resamples(list(enet =enet.fit, lasso = lasso.fit, pcr = pcr.fit, pls = pls.fit, gam=gam.fit,
summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: enet, lasso, pcr, pls, gam, tree, rf, boosting
## Number of resamples: 50
##
## MAE
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
## enet      0.204483979 0.21367310 0.22406282 0.22380451 0.23165551 0.24666446
## lasso     0.316133461 0.34329906 0.35120802 0.35226982 0.36352276 0.37732984
## pcr       0.178944265 0.18824442 0.19810593 0.19799056 0.20622345 0.21747534
## pls       0.178943285 0.18816450 0.19819321 0.19800937 0.20627703 0.21764404
## gam       0.030907527 0.03313992 0.03433917 0.03432081 0.03536590 0.04008940
## tree      2.511276986 2.74260652 2.84829024 2.88825200 3.00423418 3.68344338
## rf        0.041952710 0.07184448 0.10474823 0.12563785 0.17484122 0.31405421
## boosting  0.009534808 0.02372112 0.03227388 0.03738471 0.04544817 0.08821171
##           NA's
## enet      0
## lasso     0
## pcr       0
## pls       0
## gam       0
## tree      0
## rf        0
## boosting  0
##
## RMSE
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
## enet      0.26938541 0.29772082 0.32704779 0.32684611 0.34811402 0.38787385
## lasso     0.41847990 0.46740296 0.48034811 0.48332882 0.50435816 0.54279532
## pcr       0.23794846 0.28077381 0.31427071 0.31349502 0.33886646 0.39727401
## pls       0.23790094 0.28094331 0.31440206 0.31351276 0.33884139 0.39728622
## gam       0.03905939 0.04728467 0.05492136 0.05513242 0.06175933 0.08638066
## tree      2.95780787 3.89787450 4.65972092 5.05439196 5.50287083 11.28748560
## rf        0.18453821 0.50448945 0.69436904 0.99714001 1.30883301 3.07490317
## boosting  0.03440117 0.11885121 0.19221470 0.26713252 0.34138146 0.67141072
##           NA's
## enet      0
## lasso     0
## pcr       0
## pls       0
## gam       0
## tree      0
## rf        0
## boosting  0
##
## Rsquared
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## enet      0.6144698 0.6730869 0.7022847 0.7168417 0.7619367 0.8361868 0
## lasso     0.6072387 0.6688274 0.7144653 0.7218806 0.7724960 0.8578139 0
## pcr       0.6161576 0.6779777 0.7119874 0.7224607 0.7702780 0.8427547 0
## pls       0.6161506 0.6781600 0.7118646 0.7224473 0.7700829 0.8426681 0
## gam       0.9811190 0.9891391 0.9905767 0.9907163 0.9931621 0.9954404 0
```



```
## tree      0.9167779 0.9664709 0.9762732 0.9718544 0.9815702 0.9891014    0
## rf        0.9931318 0.9990826 0.9994931 0.9988237 0.9996636 0.9999535    0
## boosting 0.9995095 0.9998702 0.9999628 0.9998967 0.9999834 0.9999977    0
```

```
bwplot(resamp, metric = "RMSE")
```

